

Building Kafka-based Microservices with Akka Streams and Kafka Streams

Boris Lublinsky and Dean Wampler, Lightbend

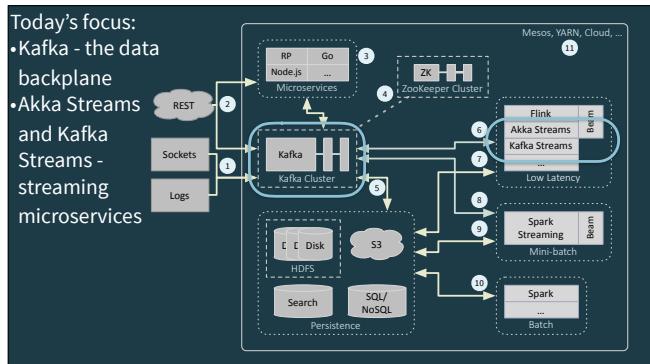
boris.lublinsky@lightbend.com
dean.wampler@lightbend.com



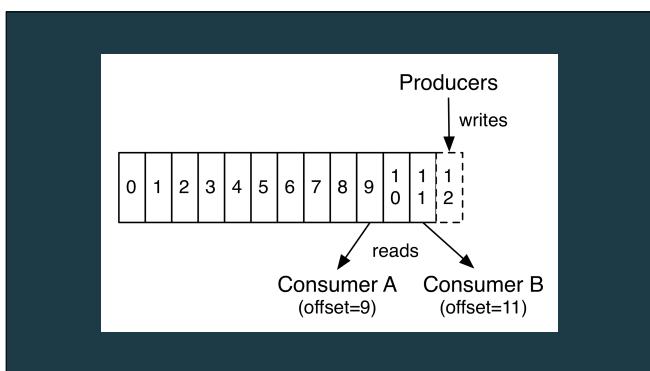
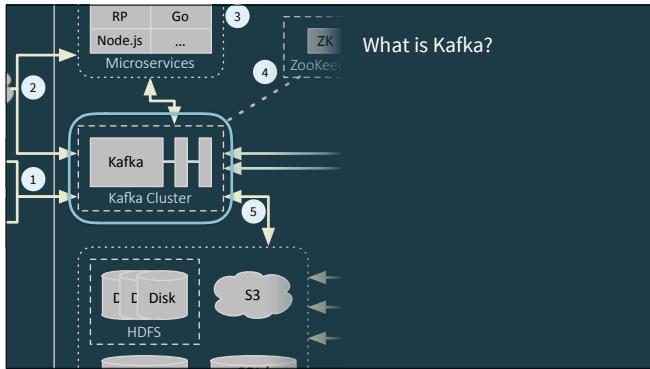
©Copyright 2018, Lightbend, Inc.
Apache 2.0 License. Please use as you see fit, but attribution is requested.



- Dean wrote this report describing the whole fast data landscape.
- bit.ly/lightbend-fast-data
- Previous talks (“Stream All the Things!”) and webinars (such as this one, <https://info.lightbend.com/webinar-moving-from-big-data-to-fast-data-here-how-to-pick-the-right-streaming-engine-recording.html>) have covered the whole architecture. This session dives into the next level of detail, using Akka Streams and Kafka Streams to build Kafka-based microservices



Kafka is the data backplane for high-volume data streams, which are organized by topics. Kafka has high scalability and resiliency, so it's an excellent integration tool between data producers and consumers.



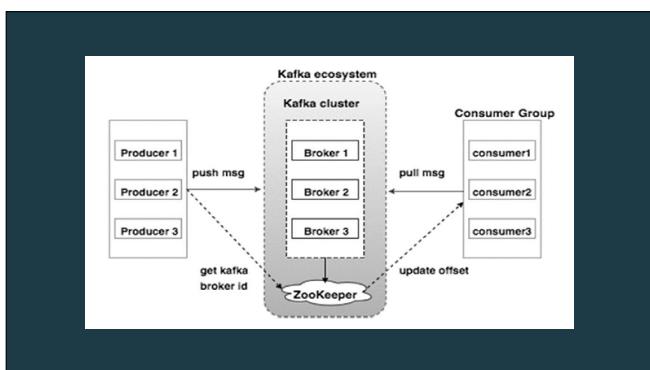
Kafka is a distributed log, storing messages sequentially. Producers always write to the end of the log, consumers can read on the log offset that they want to read from (earliest, latest, ...)

Kafka can be used as either a queue or pub sub

The main differences are:

1. Log is persistent where queue is ephemeral (reads pop elements)
2. Traditional message brokers manage consumer offsets, while log systems allow users to manage offsets themselves

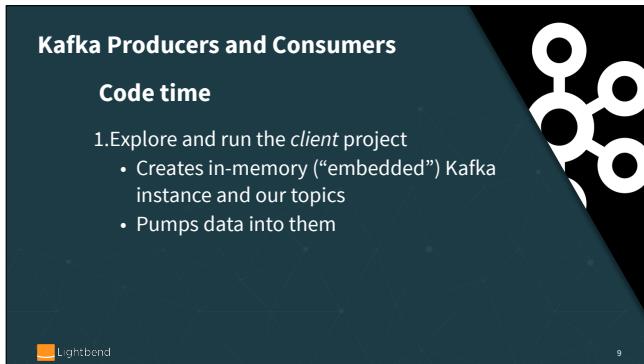
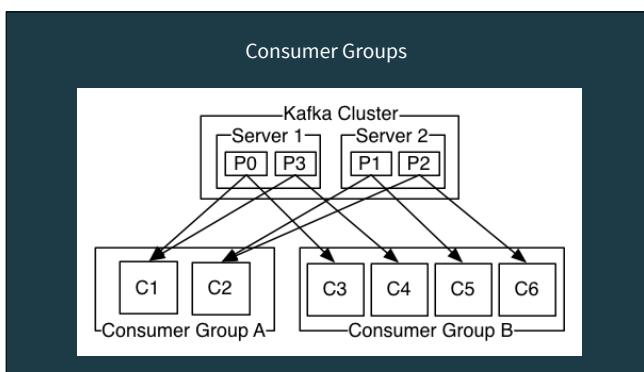
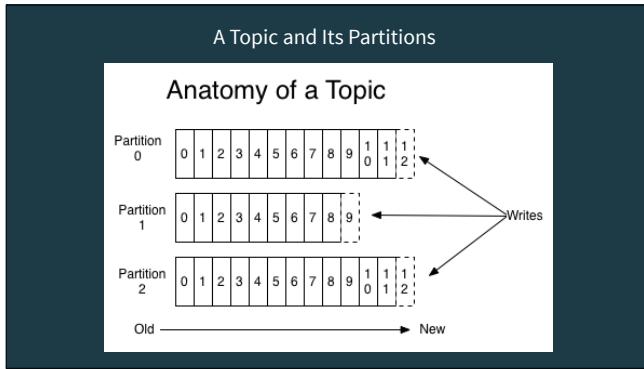
Alternatives to Kafka include Pravega (EMC) and Distributed Log/Pulsar (Apache)



Kafka cluster typically consists of multiple brokers to maintain load balance.

One Kafka broker instance can handle hundreds of thousands of reads and writes per second and each broker can handle TB (based on the disk size and network performance) of messages without performance impact. Kafka broker leader election can be done by ZooKeeper.

Image: Apache Kafka website



Kafka data is organized by topic

A topic can be comprised of multiple partitions.

A partition is a physical data storage artifact. Data in a partition can be replicated across multiple brokers. Data in a partition is guaranteed to be sequential.

So, a topic is a logical aggregation of partitions. A topic doesn’t provide any sequential guarantee (except a one-partition topic, where it’s “accidental”).

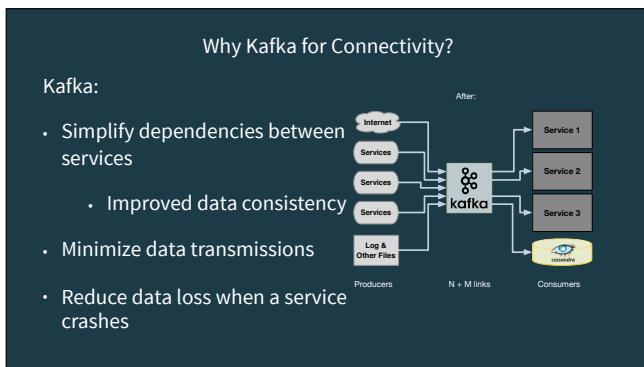
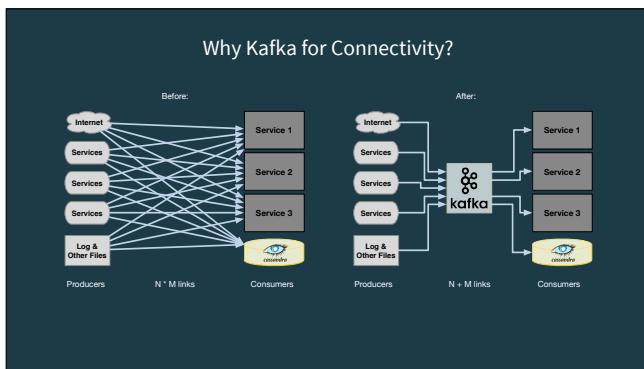
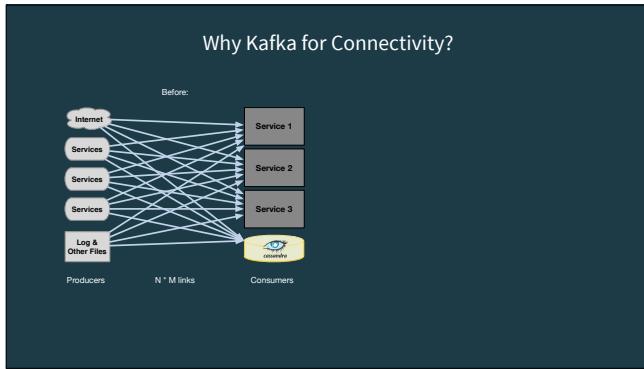
Partitioning is an important scalability mechanism - individual consumers can read dedicated partitions.

Partitioning mechanisms - round-robin, key (hash) based, custom. Consider

Consumers label themselves with a consumer group name, and each record published to a topic is delivered to one consumer instance within each subscribing consumer group (compare to queue semantics in traditional messaging). Consumer instances can be in separate processes or on separate machines.

Image: Apache Kafka website

The embedded Kafka approach is suitable for non-production scenarios only, like learning ;)



Several problems here: 1) Services are coupled! 2) What if Service 1 crashes; we might lose data from all the upstream producers connected to it. 3) Every producer-consumer pair has to understand the API and behavior of its “peer”. 4) It’s hard to understand what’s going on.

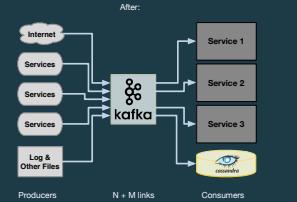
Kafka can simplify the situation by providing a single backbone which is used by all services (there are of coarse topics, but they are more logical then physical connections). Additionally Kafka persistence provides robustness when a service crashes (data is captured safely, waiting for the service to be restarted) - see also temporal decoupling, and provide the simplicity of one “API” for communicating between services.

Kafka can significantly improve decoupling (no service specific endpoints, temporal decoupling), It minimize the amount of data send over network, each producer writes data to Kafka, instead of writing it to multiple consumers. This also improves data consistency - the same data is consumed by all consumers. Extensibility is greatly simplified - adding new consumers does not require any changes to producers, and provide the simplicity of one “API” for communicating between services.

Why Kafka for Connectivity?

Kafka:

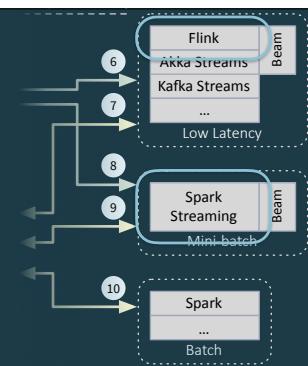
- M producers, N consumers
 - Improved extensibility
- Simplicity of one “API” for communication



Kafka can significantly improve decoupling (no service specific endpoints, temporal decoupling). It minimizes the amount of data sent over network, each producer writes data to Kafka, instead of writing it to multiple consumers. This also improves data consistency - the same data is consumed by all consumers. Extensibility is greatly simplified - adding new consumers does not require any changes to producers, and provide the simplicity of one “API” for communicating between services.

Streaming Engines:

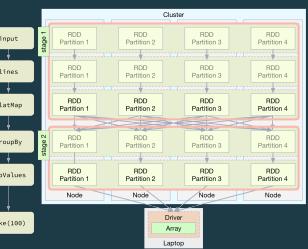
Spark, Flink - services to which you submit work. Large scale, automatic data partitioning.



They support highly scalable jobs, where they manage all the issues of scheduling processes, etc. You submit jobs to run to these running daemons. They handle scalability, failover, load balancing, etc. for you.

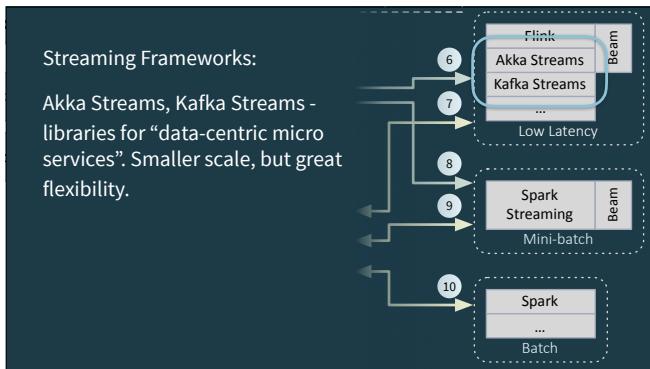
Streaming Engines:

Spark, Flink - services to which you submit work. Large scale, automatic data partitioning.



You have to write jobs, using their APIs, that conform to their programming model. But if you do, Spark and Flink do a great deal of work under the hood for you!

An example of how Spark decomposes your logical data flow or query into “stages” each of which has one JVM per data “partition”. Spark also handles partitioning for you. Flink works in a similar way.



Much more flexible deployment and configuration options, compared to Spark and Flink, but more effort is required by you to run them. They are “just libraries”, so there is a lot of flexibility and interoperation capabilities.



Scott Hanselman @shanselman [Follow](#)

Microservices, for when your in-process methods have too little latency.

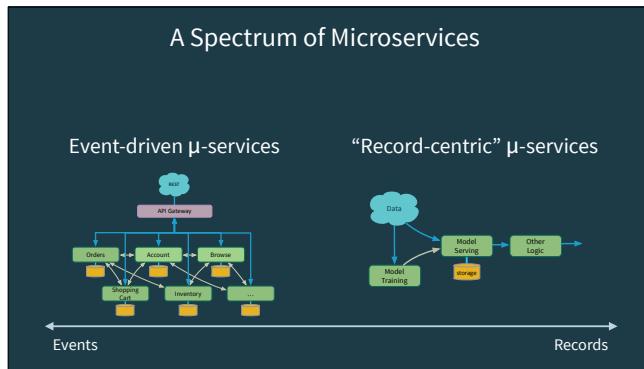
Dave Cheney @davecheney
Microservices, for when function calls are too reliable.

4:11 AM - 25 Feb 2018

207 Retweets 566 Likes

Lightbend

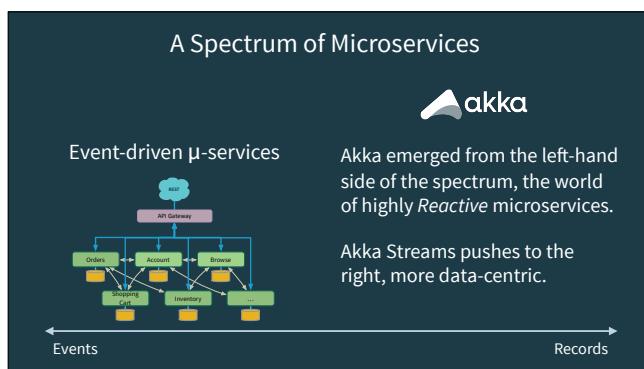
<https://twitter.com/shanselman/status/967703711492423682>



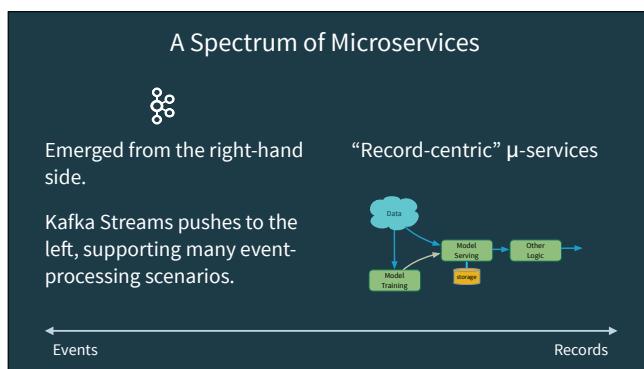
By event-driven microservices, we mean that each individual datum is treated as a specific event that triggers some activity, like steps in a shopping session. Each event requires individual handling, routing, responses, etc. REST, CQRS, and Event Sourcing are ideal for this.

Records are uniform (for a given stream), they typically represent instantiations of the same information type, for example time series; we can process them individually or as a group, for efficiency.

It's a spectrum because we might take those events and also route them through a data pipeline, like computing statistics or scoring against a machine learning model (as here), perhaps for fraud detection, recommendations, etc.



We think it's useful to reflect on the history of these toolkits, because their capabilities reflect their histories. Akka Actors emerged in the world of building *Reactive* microservices, those requiring high resiliency, scalability, responsiveness, CEP, and must be event driven. Akka is extremely lightweight and supports extreme parallelism, including across a cluster. However, the Akka Streams API is effectively a dataflow API, so it nicely supports many streaming data scenarios, allowing Akka to cover more of the spectrum than before.

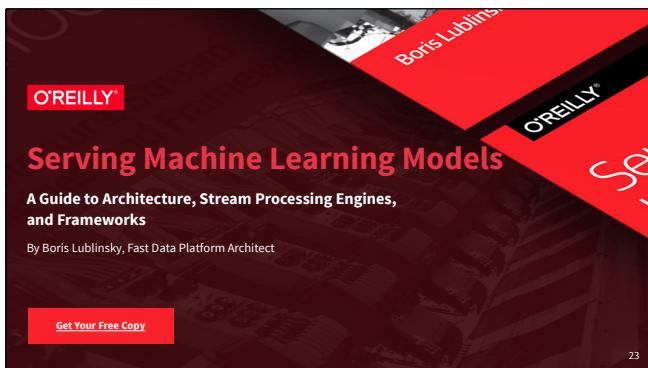


Kafka reflects the heritage of moving and managing streams of data, first at LinkedIn. But from the beginning it has been used for event-driven microservices, where the “stream” contained events, rather than records. Kafka Streams fits squarely in the record-processing world, where you define data flows for processing and even SQL. It can also be used for event processing scenarios.

Machine Learning and Model Serving: A Quick Introduction

Lightbend

We'll return to more details about AS and KS as we get into implementation details.



Our concrete examples are based on the content of this report by Boris, on different techniques for serving ML models in a streaming context.

ML Is Simple



Lightbend

Get a lot of data
Sprinkle some magic
And be happy with results

Maybe Not



Lightbend

25

Not only the climb is steep, but you are not sure which peak to climb
Court of the Patriarchs at Zion National park

Even If There Are Instructions

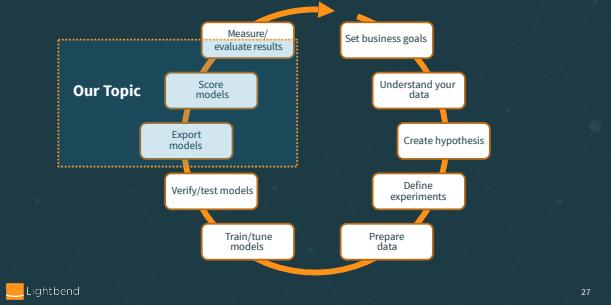


Lightbend

26

Not only the climb is steep, but you are not sure which peak to climb
Court of the Patriarchs at Zion National park

The Reality



27

- But what does a company in the, say, *Aware* stage look like, vs a company in the *Expand* stage?
- Some real-world examples can help drive that understanding.

What Is The Model?

A model is a function transforming inputs to outputs - $y = f(x)$

for example:

Linear regression: $y = a_0 + a_1 * x_1 + \dots + a_n * x_n$

Neural network: $f(x) = K(\sum_i w_i g_i(x))$

Such a definition of the model allows for an easy implementation of model's composition. From the implementation point of view it is just function composition



28

- But what does a company in the, say, *Aware* stage look like, vs a company in the *Expand* stage?
- Some real-world examples can help drive that understanding.

Model Learning Pipeline

UC Berkeley AMPLab introduced **machine learning pipelines** as a graph defining the complete chain of data transformation.



29

UC Berkeley AMPLab introduced machine learning pipelines as a graph defining the complete chain of data transformation

The advantage of such approach

It captures the whole processing pipeline including data preparation transformations, machine learning itself and any required post processing of the ML results.

Although a single predictive model is shown on this picture, in reality several models can be chained to gather or composed in any other way. See PMML documentation for description of different model composition approaches.

Definition of the complete model allows for optimization of the data

Traditional Approach to Model Serving

- Model is code
- This code has to be saved and then somehow imported into model serving

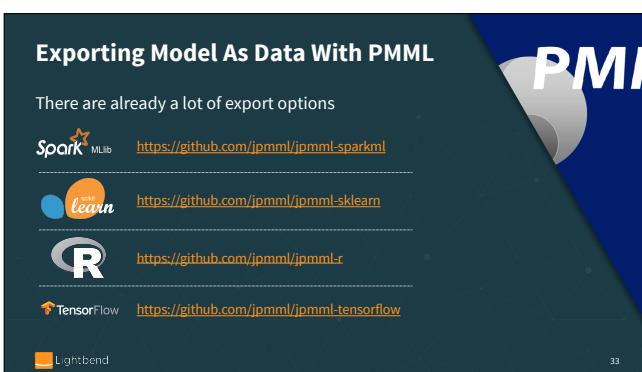
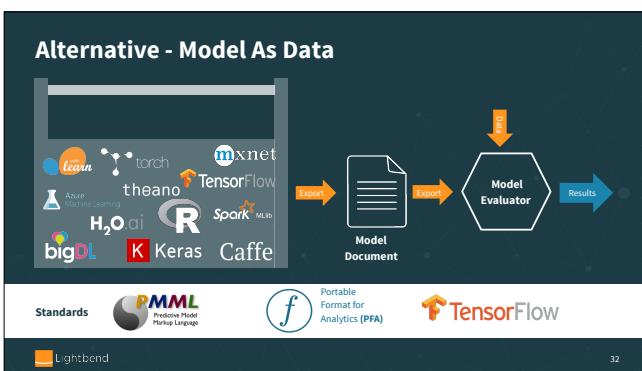
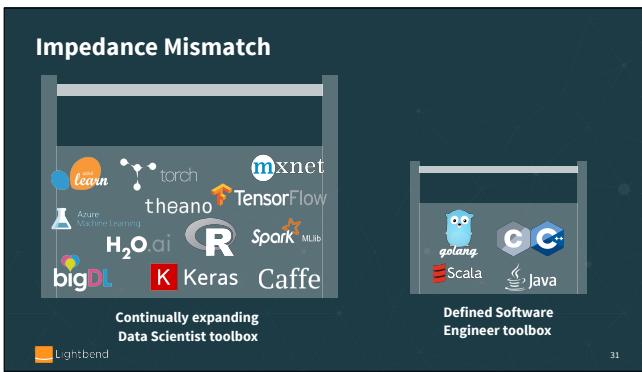
Why is this problematic?

30

This provides guidance only and should be used as a suggestion. Choose the best fit for each answer. Return to this deck once the questionnaire is complete and scored.

Link: <https://goo.gl/forms/cay1pMtxIQh83cSD3>

NOTE: Score answers on a 1 to 4 point scale for each answer from top to bottom -



In his talk at the last Flink Forward, Ted Dunning discussed the fact that with multiple tools available to Data scientists, they tend to use different tools for solving different problems and as a result they are not very keen on tools standardization. This creates a problem for software engineers trying to use “proprietary” model serving tools supporting specific machine learning technologies. As data scientists evaluate and introduce new technologies for machine learning, software engineers are forced to introduce new software packages supporting model scoring for these additional technologies.

In order to overcome these differences, Data Mining Group have introduced 2 standards - Predictive Model Markup Language (PMML) and Portable Format for Analytics (PFA), both suited for description of the models that need to be served. Introduction of these models led to creation of several software products dedicated to “generic” model serving, for example Openscoring, Open data group, etc.

Another de facto standard for machine learning is Tensorflow, which is widely used for both machine learning and model serving. Although it is a proprietary format, it is used so widely that it becomes a standard

The result of this standardization is creation of the open source projects, supporting these formats - JPMML and Hadrian which are gaining more and

Evaluating PMML Model

There are also a few PMML evaluators

 Java <https://github.com/jpmml/jpmml-evaluator>

 python <https://github.com/opendatagroup/augustus>

 Lightbend

34

Exporting Model As Data With Tensorflow

- Tensorflow execution is based on Tensors and Graphs
- Tensors are defined as multilinear functions which consist of various vector variables
- A computational graph is a series of Tensorflow operations arranged into graph of nodes
- Tensorflow supports exporting graphs in the form of binary protocol buffers
- There are two different export format - optimized graph and a new format - saved model

 Lightbend

35

Evaluating Tensorflow Model

- Tensorflow is implemented in C++ with a Python interface.
- In order to simplify Tensorflow usage from Java, in 2017 Google introduced Tensorflow Java API.
- Tensorflow Java API supports importing an exported model and allows to use it for scoring.

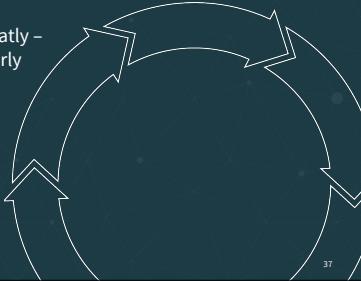
 Lightbend

36

We have a previously-trained TF model on the included “Wine Records” data. We’ll import that model to do scoring.

Additional Considerations – Model Lifecycle

- Models tend to change
- Update frequencies vary greatly – from hourly to quarterly/yearly
- Model version tracking
- Model release practices
- Model update process

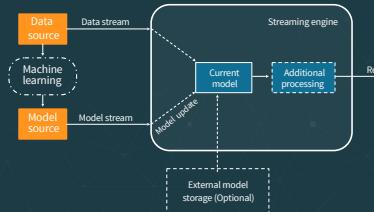


Lightbend

37

The Solution

A streaming system allowing to update models without interruption of execution (dynamically controlled stream).



Lightbend

38

The majority of machine learning implementations are based on running model serving as a REST service, which might not be appropriate for high-volume data processing or streaming systems, since they require recoding/restarting systems for model updates. For example, Flink TensorFlow or Flink JPPML.

Model Representation (Protobufs)

```
// On the wire
syntax = "proto3";
// Description of the trained model.
message ModelDescriptor {
    string name = 1; // Model name
    string description = 2; // Human readable
    string dataType = 3; // Data type for which this model is applied.
    enum ModelType { // Model type
        TENSORFLOW = 0;
        TENSORFLOWSAVED = 2;
        PMML = 2;
    };
}
message MessageContent {
    ModelType modeltype = 4;
    oneof MessageContent {
        bytes data = 5;
        string location = 6;
    }
}
```

Lightbend

39

You need a neutral representation format that can be shared between different tools and over the wire. Protobufs (from Google) is one of the popular options. Recall that this is the format used for model export by TensorFlow. Here is an example.

Model Representation (Scala)

```
trait Model {  
    def score(input: AnyVal) : AnyVal  
    def cleanup() : Unit  
    def toBytes() : Array[Byte]  
    def getType : Long  
}  
  
def ModelFactory{  
    def create(input: ModelDescriptor) : Model  
    def restore(bytes : Array[Byte]) : Model  
}
```



40

Corresponding Scala code that can be generated from the description.

Side Note: Monitoring

Model monitoring should provide information about usage, behavior, performance and lifecycle of the deployed models

```
case class ModelToServeStats(  
    name: String, // Model name  
    description: String, // Model descriptor  
    modelType: ModelDescriptor.ModelType, // Model type  
    since: Long, // Start time of model usage  
    var usage: Long = 0, // Number of servings  
    var duration: Double = 0.0, // Time spent on serving  
    var min: Long = Long.MaxValue, // Min serving time  
    var max: Long = Long.MinValue // Max serving time  
)
```

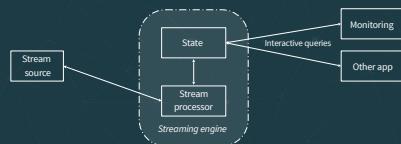


41

Queryable State

Queryable state: ad hoc query of the state in the stream. Different than the normal data flow.

Treats the stream processing layer as a lightweight embedded database. Directly query the current state of a stream processing application. No need to materialize that state to a database, etc. first.



42

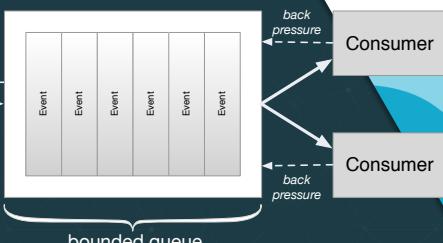
Kafka Streams and Flink have built-in support for this and its being added to Spark Streaming. We'll show how to use other Akka features to provide the same ability in a straightforward way for Akka Streams.

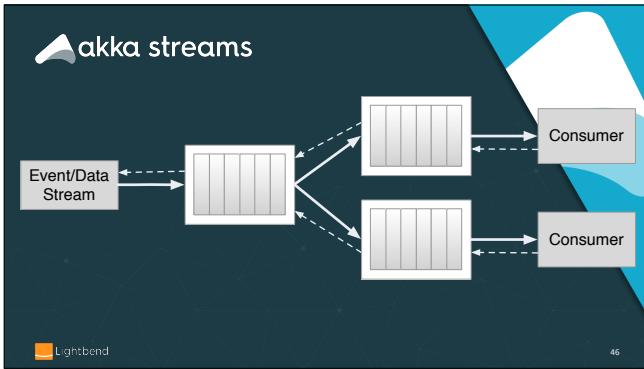
- We'll work with Akka Streams examples first

- A library
- Implements Reactive Streams.
- <http://www.reactive-streams.org/>
- Back pressure for flow control

See this website for details on why *back pressure* is an important concept for reliable flow control, especially if you don't use something like Kafka as your "near-infinite" buffer between services.

Bounded queues are the only sensible option (even Kafka topic partitions are bounded by disk sizes), but to prevent having to drop input when it's full, consumers signal to producers to limit flow. Most implementations use a push model when flow is fine and switch to a pull model when flow control is needed.

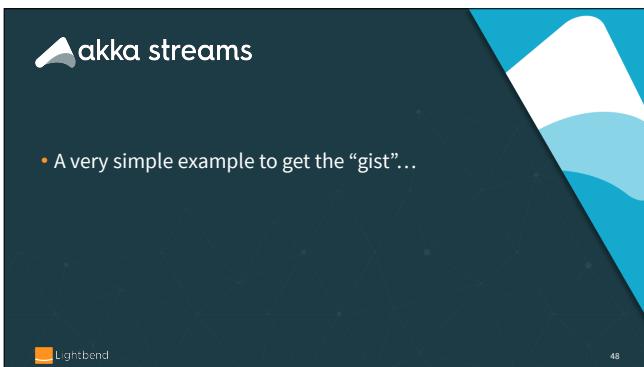




And they compose so you get end-to-end back pressure.



Rich, mature tools for the full spectrum of microservice development. Akka Streams adds a streaming abstraction on top of Akka actors. It's ideal when complex event processing (CEP) is the preferred model, as opposed to in bulk processing of data. Akka's powerful Actor model abstracts over the details of thread programming for highly concurrent apps, with libraries for clustering, persisting state, and data interchange with many sources and sinks (the "Alpakka" project).



```
import akka.stream._  
import akka.stream.scaladsl._  
import akka.{ NotUsed, Done }  
import akka.actor.ActorSystem  
import scala.concurrent._  
import scala.concurrent.duration._  
  
implicit val system = ActorSystem("QuickStart")  
implicit val materializer = ActorMaterializer()  
  
val source: Source[Int, NotUsed] = Source(1 to 10)  
val factorials = source.scan(BigInt(1)) ( (acc, next) => acc * next )  
factorials.runWith(Sink.foreach(println))
```

This example is in akkaStreamsCustomStage/simple-akka-streams-example.sc

```
import akka.stream._  
import akka.stream.scaladsl._  
import akka.{ NotUsed, Done }  
import akka.actor.ActorSystem  
import scala.concurrent._  
import scala.concurrent.duration._  
  
implicit val system = ActorSystem("QuickStart")  
implicit val materializer = ActorMaterializer()  
  
val source: Source[Int, NotUsed] = Source(1 to 10)  
val factorials = source.scan(BigInt(1)) ( (acc, next) => acc * next )  
factorials.runWith(Sink.foreach(println))
```

This example is in akkaStreamsCustomStage/simple-akka-streams-example.sc

```
import akka.stream._  
import akka.stream.scaladsl._  
import akka.{ NotUsed, Done }  
import akka.actor.ActorSystem  
import scala.concurrent._  
import scala.concurrent.duration._  
  
implicit val system = ActorSystem("QuickStart")  
implicit val materializer = ActorMaterializer()  
  
val source: Source[Int, NotUsed] = Source(1 to 10)  
val factorials = source.scan(BigInt(1)) ( (acc, next) => acc * next )  
factorials.runWith(Sink.foreach(println))
```

This example is in akkaStreamsCustomStage/simple-akka-streams-example.sc

```
import akka.stream._  
import akka.stream.scaladsl._  
import akka.{ NotUsed, Done }  
import akka.actor.ActorSystem  
import scala.concurrent._  
import scala.concurrent.duration._  
  
implicit val system = ActorSystem("QuickStart")  
implicit val materializer = ActorMaterializer()  
  
val source: Source[Int, NotUsed] = Source(1 to 10)  
val factorials = source.scan(BigInt(1)) ( (acc, next) => acc * next )  
factorials.runWith(Sink.foreach(println))
```

Create a Source.
Scan it and compute factorials, output to a Sink, and run it.

This example is in akkaStreamsCustomStage/simple-akka-streams-example.sc

```
import akka.stream._  
import akka.stream.scaladsl._  
import akka.{ NotUsed, Done }  
import akka.actor.ActorSystem  
import scala.concurrent._  
import scala.concurrent.duration._  
  
implicit val system = ActorSystem("QuickStart")  
implicit val materializer = ActorMaterializer()  
  
val source: Source[Int, NotUsed] = Source(1 to 10)  
val factorials = source.scan(BigInt(1)) ( (acc, next) => acc * next )  
factorials.runWith(Sink.foreach(println))
```

Create a Source.
Scan it and compute factorials, output to a Sink, and run it.

The core concepts are sources and sinks, connected by flows. There is the notion of a Graph for more complex dataflows, but we won't discuss them further

akka streams

- This example is included in the project:
 - akkaStreamsCustomStage/simple-akka-streams-example.sc
- To run it (showing the different prompt!):

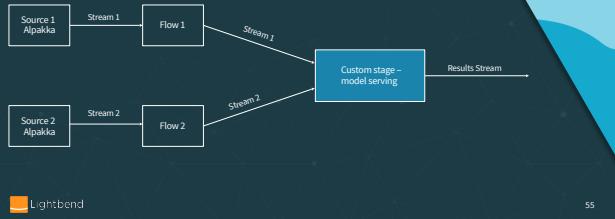
```
$ sbt  
sbt:akkakafkaTutorial> project akkaStreamsCustomStage  
sbt:akkakafkaTutorial> console  
scala> :load akkaStreamsCustomStage/simple-akka-streams-example.sc
```

Lightbend

The “.sc” extension is used so that the compiler doesn’t attempt to compile this Scala “script”. Using “.sc” is an informal convention for such files.

Using Custom Stage

Create a custom stage, a fully type-safe way to encapsulate new functionality. Like adding a new “operator”.



Custom stage is an elegant implementation but doesn't scale well to a large number of models. Although a stage can contain a hash map of models, all of the execution will be happening at the same place

Using a Custom Stage

Code time

1. Walk through the whole tutorial project
2. Run the *client* project (if not already running)
3. Explore and run *akkaStreamsCustomStage* project

Custom stage is an elegant implementation but not scale well to a large number of models. Although a stage can contain a hash map of models, all of the execution will be happening at the same place

Exercises!

We've prepared some exercises. We may not have time during the course to work on them, but take a look at the *exercise* branch in the Git project (or the separate X.Y.Z_*exercise* download).

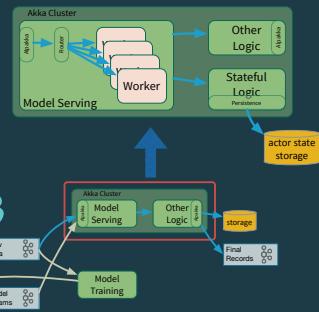
To find them, search for “// Exercise”. The *master* branch implements the solutions.

Other Production Concerns

Lightbend

58

- Scale scoring with workers and routers, across a cluster
- Persist actor state with Akka Persistence
- Connect to *almost* anything with Alpakka
- New *StreamRefs* for distributed streams
- *Lightbend Enterprise Suite*
 - for production monitoring, etc.



Here's our streaming microservice example adapted for Akka Streams. We'll still use Kafka topics in some places and assume we're using the same implementation for the "Model Training" microservice. Alpakka provides the interface to Kafka, DBs, file systems, etc. We're showing two microservices as before, but this time running in Akka Cluster, with direct messaging between them. We'll explore this a bit more after looking at the example code.

StreamRefs were just released (Akka 2.5.10). See [this video](#).

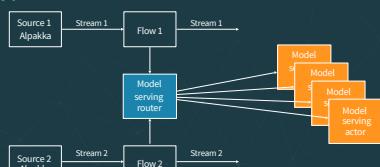
Improve Scalability for Model Serving

Use a router actor to forward requests to the actor responsible for processing requests for a specific model type.

Lightbend

60

We here create a routing layer: an actor that will implement model serving for specific model (based on key) and route messages appropriately. This way our system will serve models in parallel.



Akka Streams with Actors and Persistence

Code time

1. While still running the *client* project...
2. Explore and run *akkaActorsPersistent* project

Lightbend

61

Custom stage is an elegant implementation but not scale well to a large number of models. Although a stage can contain a hash map of models, all of the execution will be happening at the same place

More Production Concerns

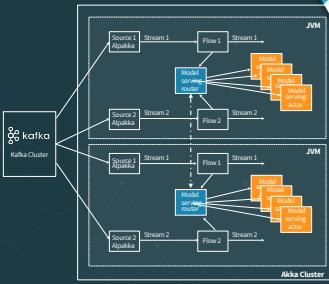
Lightbend

62

Using Akka Cluster

Two levels of scalability:

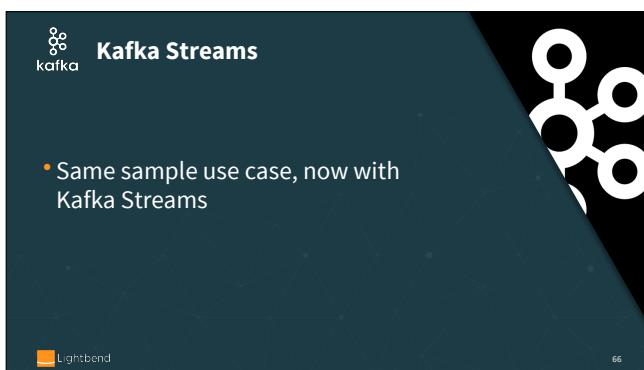
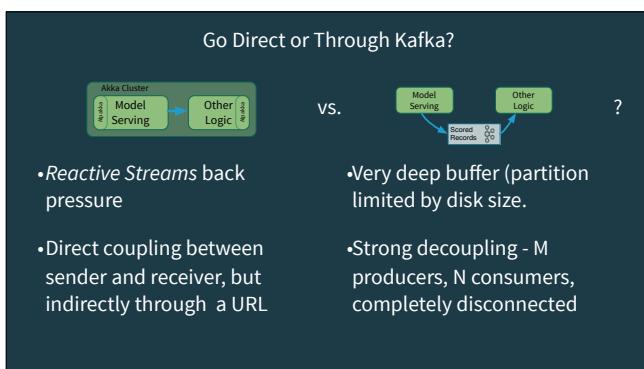
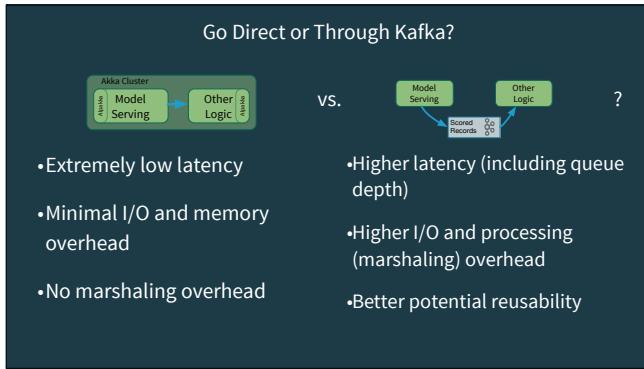
- Kafka partitioned topic allow to scale listeners according to the amount of partitions.
- Akka cluster sharing allows to split model serving actors across clusters.



Lightbend

63

A great article <http://michalplachta.com/2016/01/23/scalability-using-sharding-from-akka-cluster/> goes into a lot of details on both implementation and testing



Design choice: When is it better to use direct actor-to-actor (or service-to-service) messaging vs. going through a Kafka topic?

Design choice: When is it better to use direct actor-to-actor (or service-to-service) messaging vs. going through a Kafka topic?

Kafka Streams

- Important stream-processing concepts, e.g.,
 - Distinguish between *event time* and *processing time*
 - Windowing support
 - For more on these concepts, see
 - Dean's book ;)
 - Talks, blog posts, writing by Tyler Akidau

 Lightbend



There's a maturing body of thought about what streaming semantics should be, too much to discuss here. Dean's book provides the next level of details. See Tyler's work (from the Google Apache Beam team) for deep dives.

Kafka Streams

- KStream - per-record transformations
- KTable - last value per key
- Efficient management of application state

 Lightbend



There is a duality between streams and tables. Tables are the latest state snapshot, while streams record the history of state evolution. A common way to implement databases is to use an event (or change) log, then update the state from the log.

Kafka Streams

- Low overhead
- Read from and write to Kafka topics, memory
 - Could use Kafka Connect for other sources and sinks
- Load balance and scale based on partitioning of topics
- Built-in support for Queryable State

 Lightbend



Kafka Streams

- Two types of APIs:
- Process Topology
- Compare to [Apache Storm](#)
- DSL based on collection transformations
- Compare to Spark, Flink, Scala collections.

 Lightbend

70

Kafka Streams

- Provides a Java API
- Lightbend donating a Scala API to Apache Kafka
 - <https://github.com/lightbend/kafka-streams-scala>
 - See also our convenience tools for distributed, queryable state: <https://github.com/lightbend/kafka-streams-query>
- SQL!

 Lightbend

71

The `kafka-streams-query` uses a KS API to find all the partitions across a cluster for a given topic, query their state, and aggregate the results, behind a web service. Otherwise, you have to query the partitions individually yourself.

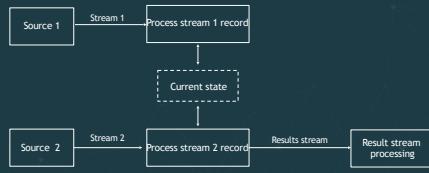
Kafka Streams

- Ideally suited for:
 - ETL -> KStreams
 - State -> KTable ???
 - Joins, including Stream and Table joins
 - “Effectively once” semantics
- Commercial support from Confluent, Lightbend, and others

 Lightbend

72

Model Serving With Kafka Streams



Lightbend

73

State Store Options We'll Explore

- “Naive”, in memory store
- Built-in key/value store provided by Kafka Streams
- Custom store

Lightbend

74

We provide three example implementations, using three different ways of storing state. “Naive” - because in-memory state is lost if the process crashes; a restart can’t pick up where the previous instance left off.

Model Serving With Kafka Streams

Code time

1. Still running the *client* project...
2. Explore and run:
`kafkaStreamsModelServerInMemoryStore`

Lightbend

75

Model Serving With Kafka Streams, KV Store

Code time (as time permits)

1. Still running the *client* project...
2. Explore and run:
`kafkaStreamsModelServerKVStore`

We probably won't get to this implementation and the next, but you can look at them on your own.

Model Serving With Kafka Streams, KV Store

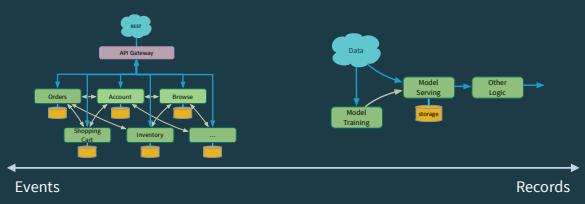
Code time (as time permits)

1. Still running the *client* project...
2. Explore and run:
`kafkaStreamsModelServerCustomStore`

To Wrap Up



Event-driven µ-services



"Record-centric" µ-services

Akka Streams is a great choice if you are building full-spectrum microservices and you need lots of flexibility in your app architectures, connecting to different kinds of data sources and sinks, etc.

Kafka Streams is a great choice if your use cases fit nicely in its “sweet spot”, you want SQL access, and you don’t need the full flexibility of something like Akka. Of course, you can use both! They are “just libraries”.

Thank You

Questions?

- Kafka streaming applications with Akka Streams and Kafka Streams (Dean)
 - Thursday 11:00 - 11:40, Expo Hall 1
- Meet the Expert (Dean)
 - Thursday 11:50 - 12:30, O'Reilly Booth, Expo Hall
- AMA, (Boris and Dean)
 - Thursday 2:40 - 3:20, 212 A-B

- And don't miss:
- Approximation data structures in streaming data processing (Debasish Ghosh)
 - Wednesday 1:50 - 2:30, 230A
 - Machine-learned model quality monitoring in fast data and streaming applications (Emre Velipasaoglu)
 - Thursday 1:50 - 2:30, LI 21 C/D

<https://www.lightbend.com/products/fast-data-platform>
boris.lublinsky@lightbend.com
dean.wampler@lightbend.com



Thank you! Please check out the other Strata San Jose sessions by Boris, Dean, and our colleagues Debasish and Emre. Check out our Fast Data Platform for commercial options for building and running microservices with Kafka, Akka Streams, and Kafka Streams.